

08/8/23

UNIT-IV

Transaction Concept

Syllabus

Transaction state, Implementation of atomicity and durability, Concurrent Executions, Serializability, Recoverability, Implementation of isolation, testing for serializability, lock-based protocol, timestamp based protocols, validation based protocols, Multiple Granularity, recovery & Atomicity, log-based recovery, Recovery with Concurrent Transactions.

Transaction state

Transaction is a set of operations which are logically related.

The main operations in transactions are

① Read operation

- It reads the data from the database and then stores it in buffer from Main memory.

Read(A) - It will read the value of 'A' from the database.

② Write operation

It writes the updated data value back.

to the database from buffer.

write (A) - will write the updated value of A from the buffer to database.

transaction properties

DBMS must ensure to maintain data in case of concurrent access and system failures. They are ACID properties.

Atomicity (All or nothing)

A transaction is said to be atomic if a transaction always executes all its actions in one step or does not execute at all.

Consistency (No violation of integrity constraints)

A transaction must preserve the consistency of db after the execution. It refers to the correctness of db.

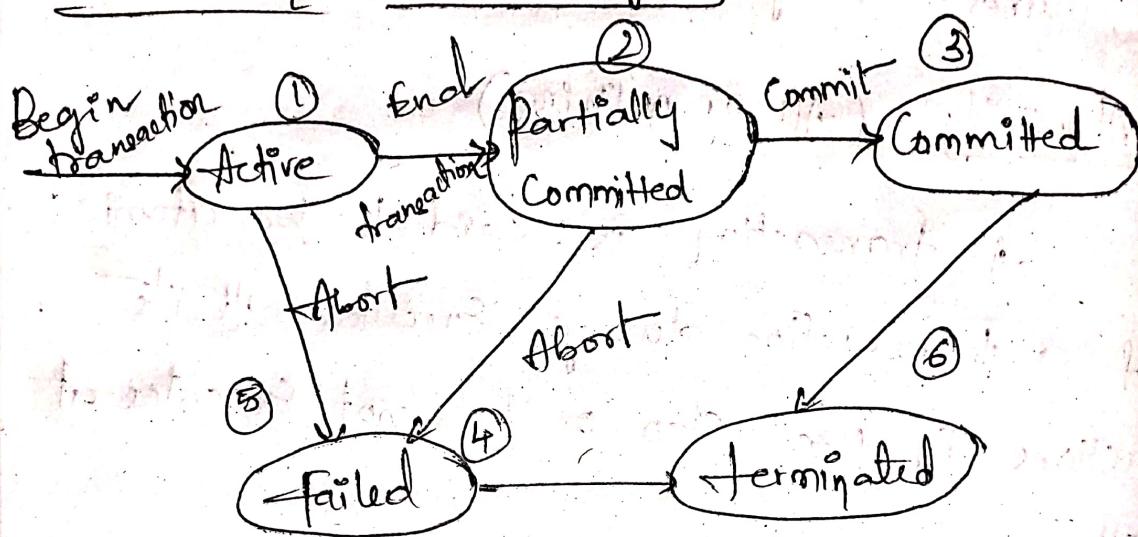
Isolation (Concurrent changes invisible)

This property ensures that multiple transactions can occur concurrently without leading to inconsistency of db. Transactions occur independently without interference.

Durability - (Committed Update Persist)

The effect of completed or committed transactions should persist even after a crash. Once a transaction commits the system must guarantee that the result of its operations.

Transaction State Diagram



Active

If a transaction enters into an active state when the execution process begins. During this state read or write operations can be performed.

Partially Committed

If a transaction goes into partially committed state after the end of a transaction.

Changes are stored in buffer but not updated in db.

Committed -

When the transaction is Committed to state, it has already completed its execution successfully. All of its changes are recorded to the db permanently.

Failed

A transaction Considered failed when any one of the checks fail(s) or if the transaction is aborted when it is in active state.

Terminated

State of transaction reaches terminated state when certain transactions are leaving the system can't be restarted.

Concurrent Execution (Multiprogramming)

Multiple transactions run Concurrently.
Multiple transactions causes several complications with consistency of data.

two good reasons for Concurrency.

- ① Improved throughput and resource utilization

Throughput is no. of transactions executed in a given amount of time.

A transaction consists of many steps. Some involve I/O activity and others involve CPU activity. The CPU and the disks in computer system can operate parallelly.

The parallelism of the CPU and the I/O system can therefore be exploited to run multiple transactions in parallel.

If one transaction is reading or writing data on disk another can be running in CPU. All of this increase throughput of the system (concur) correspondingly.

Reduced waiting time

If transaction run serially a short transaction may have to wait for a preceding long transaction to complete.

If the transactions are operating on different parts of db it is better to run them concurrently.

Serializability -

when multiple transactions are running concurrently then there is a possibility that the db may be left in inconsistent state. Serializability ^{helps us to} checks that which schedule are serializable.

It means Executing transactions one after the other.

Schedule

A seq of operations by a set of concurrent transactions that preserves the order of operations in each individual transaction.

Multiple transaction

↳ set of actions

↳ which Schedule

↳ which transaction

↳ which time.

Serial schedule

A schedule where the operations of each transaction are executed consecutively without any interleaved operations from another transaction.
 ↓
 No overlapping
 No interleaving

Non-Serial Schedule

A schedule - the operations from a set of concurrent transactions are interleaved.

Serializable Schedule

The objective of serializability is to find non-serial schedules that allow transactions to execute concurrently without interfering one another and thereby produce a db that is produced by serial execution.

If the result is same when it executes serial schedule and non-serial schedule then we can say that transaction is serializable.

Example

Let T_1 and T_2 are two transactions.
 T_1 transfers 50/- from Account 'A' to Account 'B'

$T_1: \text{read}(A);$

$A := A - 50;$

$\text{write}(A); \quad \text{read}(B);$

$B := B + 50;$

$\text{write}(B);$

T_2 transfer 10 percent of the balance from Acc 'A' to Acc 'B'

T_2 : read(A)

temp := A * 0.1

A := A - temp;

write(A);

read(B);

B := B + temp;

write(B);

Suppose two transactions are executed in the order T_1 followed by T_2 .

→ It is divided into two types.

Here all transactions should be arranged in a particular order; even if all the transaction is concurrent, and is not serializable, then it produce incorrect result.

① Conflict Serializability

It is one of the types of serializability which can be used to check whether a non-serial schedule is conflict serializable or not.

Conflicting operations

Two operations are said to be in conflict, if they satisfy the below three conditions.

e.g. the order table and the customer table

- One customer can have multiple orders, but each order only belongs to one customer.

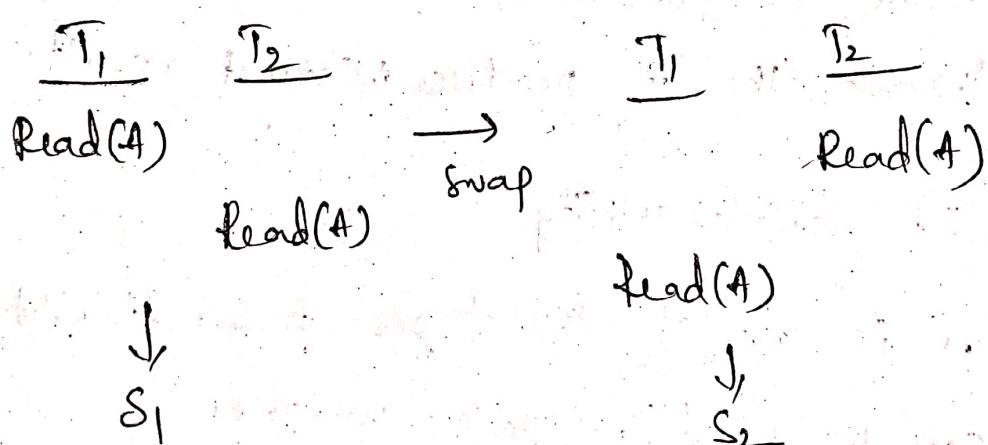
→ Both operations should have different transaction $T_1 \neq T_2$

→ Both transactions should have the same data item (x)

- There should be at least one write operation b/w the two operations. $w(x)$

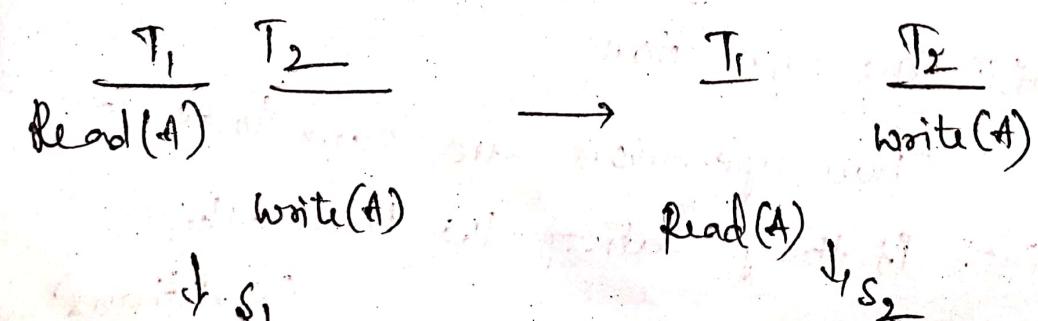
e.g. Swapping is possible only if s_1 and s_2 are logically equal.

① T_1 : Read(A) T_2 : Read(A)



$$S_1 = S_2 - \text{no-conflict}$$

② T_1 : Read(A) T_2 : Write(A)



$S_1 \neq S_2 \rightarrow \text{Conflict}$

Conflict Equivalent

Two schedules are said to be conflict equivalent if and only if:

- ① they contain the same set of transaction
- ② if each pair of conflict operations are ordered in the same way.

Non-Serial Schedule

<u>T₁</u>	<u>T₂</u>
read(A)	
write(A)	
	Read(A)
	Write(A)
read(B)	
write(B)	
	Read(A)
	Write(B)

Serial schedule

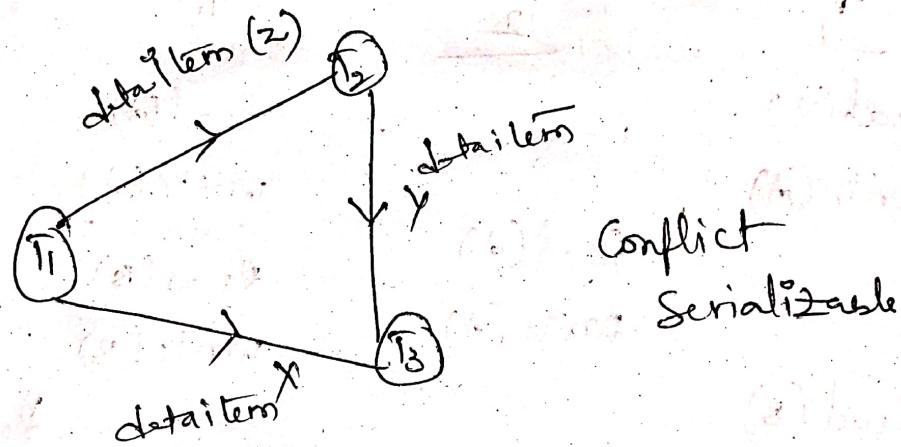
<u>T₁</u>	<u>T₂</u>
read(A)	
write(A)	
	Read(B)
	Write(B)
read(B)	
write(B)	
	Read(A)
	Write(A)
	Read(B)
	Write(B)

After swapping of non-conflict operations,
the schedule S_1 becomes:

<u>T₁</u>	<u>T₂</u>
	→ write Serial Schedule Example Here

check for conflict

<u>T₁</u>	<u>T₂</u>	<u>T₃</u>	nodes T ₁ , T ₂ , T ₃
R(x)			T ₁ → R(x)
	R(z)		T ₃ → w(x)
R(z)	R(x)		T ₂ → R(y)
	R(y)		T ₁ → R(z)
	w(z)	w(x)	T ₂ → w(z)
	w(x)	w(y)	



View Serializability

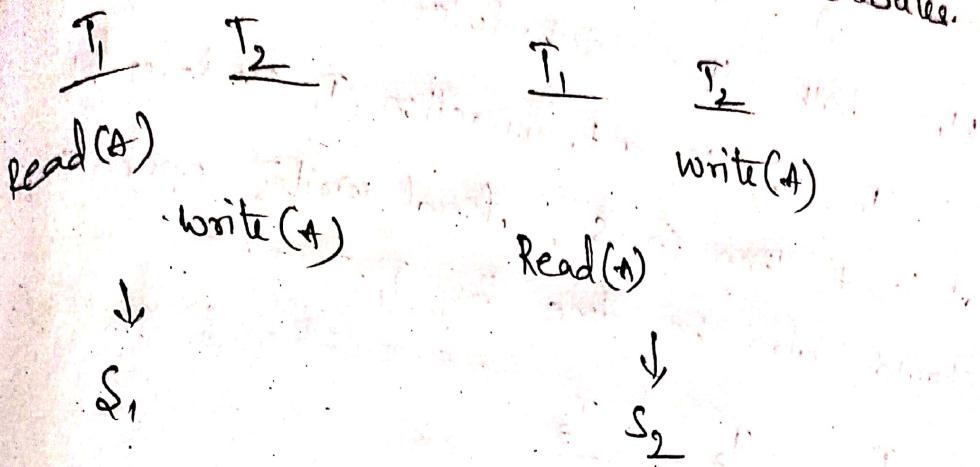
It is a process to find out the given schedule is Serializable or not.

View Equivalent

Two schedules T₁ & T₂ are said to be view equivalent if they satisfy the below conditions.

① Initial Read

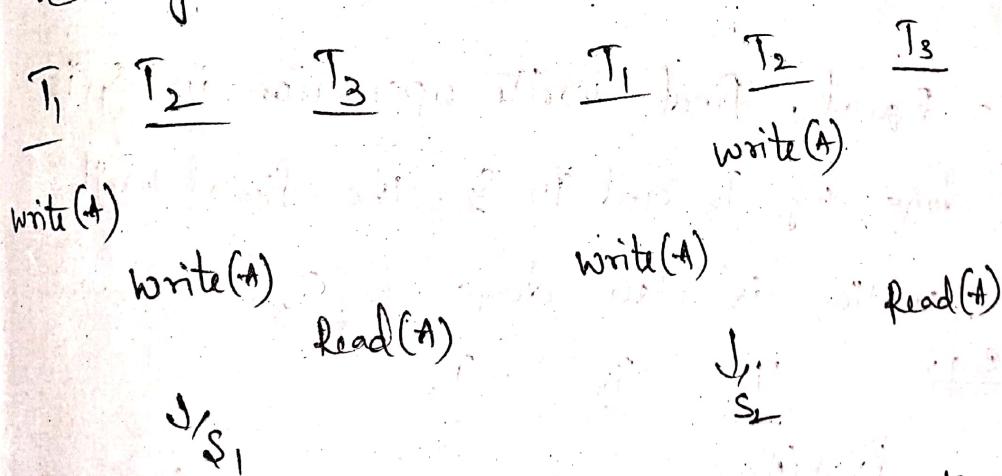
Initial Read of each data item in transactions must match in both schedules.



The two schedules are equivalent because initial read option in S_1 is done by T_1 and in S_2 it is also done by T_1 .

② Updated Read

In schedule S_1 the transaction T_1 is reading data item updated by T_2 .



+ Not view equivalent, in S_1 , T_3 is reading a update by T_2 and by S_2 , T_3 is

reading & updated by T_1 .

③ Final write

A final write must be the same b/w both the schedules.

In S_1 , if a transaction T_1 updates 'A' at least then in S_2 , final writes operations should also be done by T_1 .

T_1 T_2 T_3

write(A)

Read(A)

write(A)

$\rightarrow S_1$

T_1 T_2 T_3

Read(A)

$\rightarrow S_2$

write(A)

write(A)

- Equal, final write operation in S_1 is done by T_3 and in S_2 , the final write operation is also done by T_3 .

Eg:-

T_1 T_2 T_3

Read(A)

write(A)

$\rightarrow S$

write(A)

write(A)

→ with 3 transactions, the total no. of possible schedule $3! = 6$

$$S_1 = T_1 \ T_2 \ T_3 \quad S_3 = T_2 \ T_3 \ T_1$$

$$S_2 = T_1 \ T_3 \ T_2 \quad S_4 = T_2 \ T_1 \ T_3$$

$$S_5 = T_3 \ T_1 \ T_2$$

$$S_6 = T_3 \ T_2 \ T_1$$

* Taking first schedule S_1 :

T_1 T_2 T_3

read(A)

write(A)

write(A)

write(A)

Step 1: final update

In S_1 & S_1 , there is no read except the initial read.

Step 2: Initial Read

Condition satisfied

Step 3: final write

Condition satisfied.

$T_1 \rightarrow T_2 \rightarrow T_3$

Recoverability of Schedule

Sometimes a transaction may not execute completely due to a software issue, system crash or hardware failure.

In this case, the failed transaction has to be roll back.

But some other transaction may also have used value produced by the failed transaction. So we also have to rollback those transactions.

① Irrecoverable Schedules

The schedule will be irrecoverable if T_j reads the updated value of T_i and T_j committed before T_i committed.

Concurrency Control protocols

① Lock-Based protocols

A lock guarantees exclusive use of data items to a current transaction.

- To access data items (lock acquire)
- After completion of transaction (release lock)

lock Compatibility Matrix

	S	X
S	✓	X
X	X	X

A transaction may be granted a lock on an item if the requested lock is compatible with locks already held on the item by other transactions.

- Any ~~single~~ no. of transactions can hold shared locks on a item, but if any tran holds an Exclusive(X) on the item no other tran may hold any lock on the item.

upgrade / Downgrade locks:

A transaction that holds a lock on an item 'A' is allowed certain condition to change the lock state from one state to another

upgrade - A S(A) can be upgraded to X(A)

if T_i is the only tran holding the lock

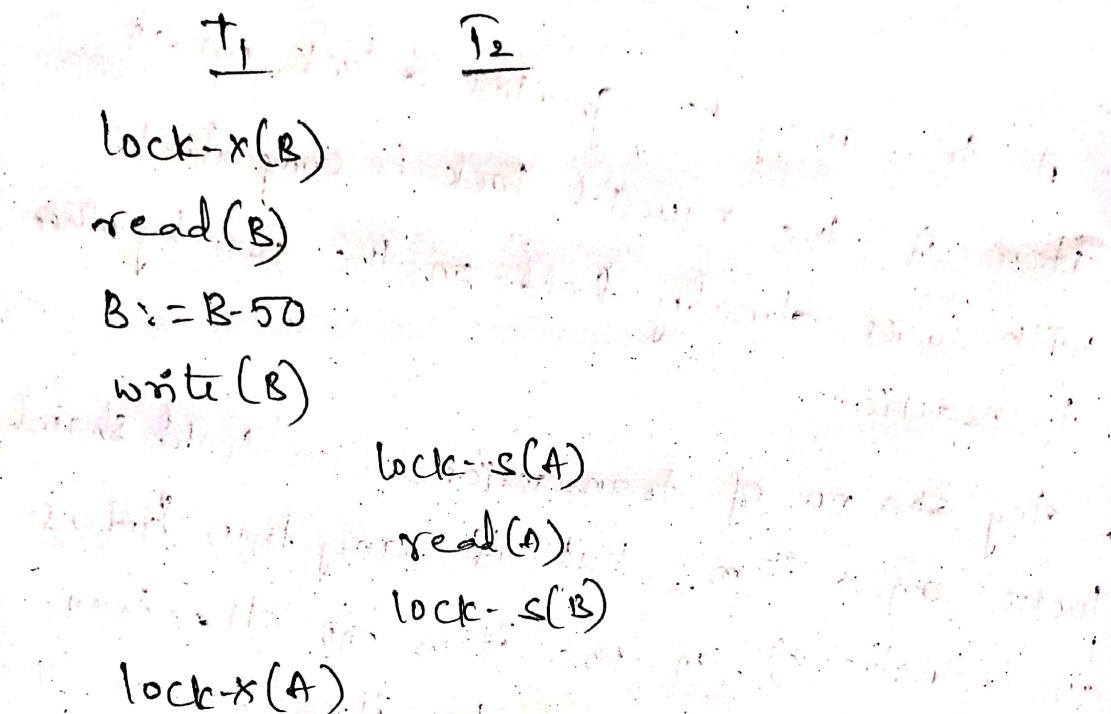
on element A

Downgrade - we may downgrade X(A) to S(A)

when we feel that we no longer want to write on data item A.

As we holding 'X-lock' on A, we need not check any conditions.

Problem with simple lock



Deadlock - T_1 holds $X(A)$ - over B,

T_2 holds $S(B)$ - over A

T_2 requests for lock on B

T_1 requests for lock on A

It imposes deadlock.

Time-stamp based protocol

- It helps DBMS to identify the transactions
- It is a unique identifier. Each transaction is issued a timestamp

they are usually assigned by order to which they are submitted to the system.

If refers as Transaction T as $TS(T)$.

old transaction $\leftarrow TS(1) \leq TS(2) \rightarrow$ new transaction.

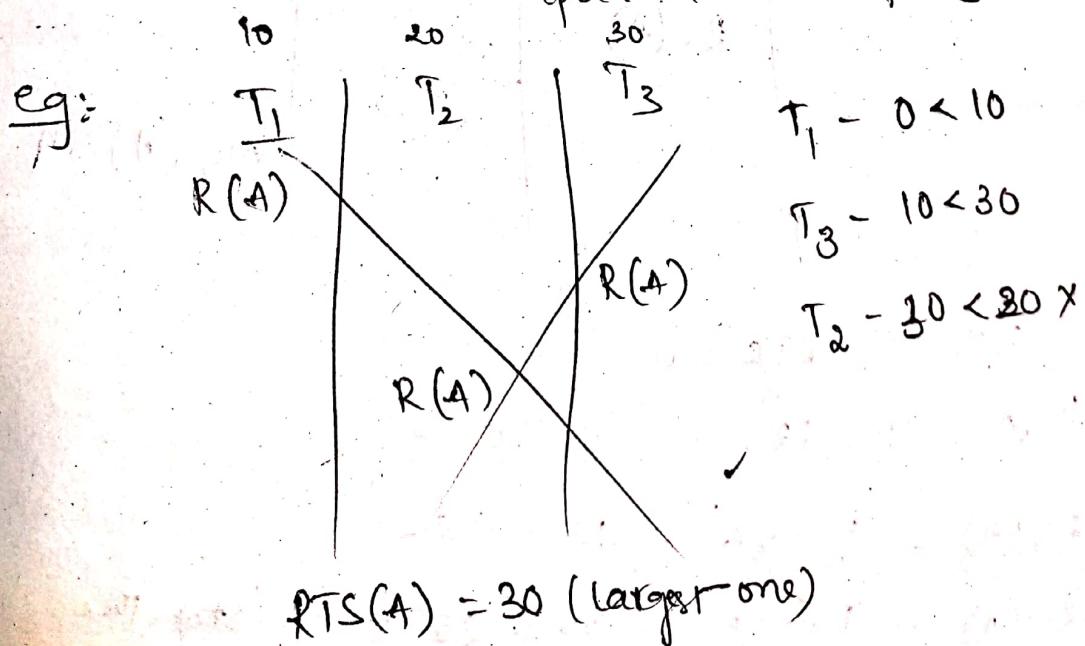
- It ensures that any conflicting read & write operations are executed in timestamp order.

The Basic TO protocol in the following two cases: (Thomas Write rules)

① T_1 issues Read(A) -

i) If $TS(T_1) < w(A)$ { T_1 is an older transaction
then = the last tran that wrote the value of A } repeat fails

ii) if $TS(T_1) \geq w(TS(A))$ { T_1 allowed to read updated value of A }



- This protocol uses system clock (or) logical clock to execute higher priority transaction first.
- As soon as transaction is created,

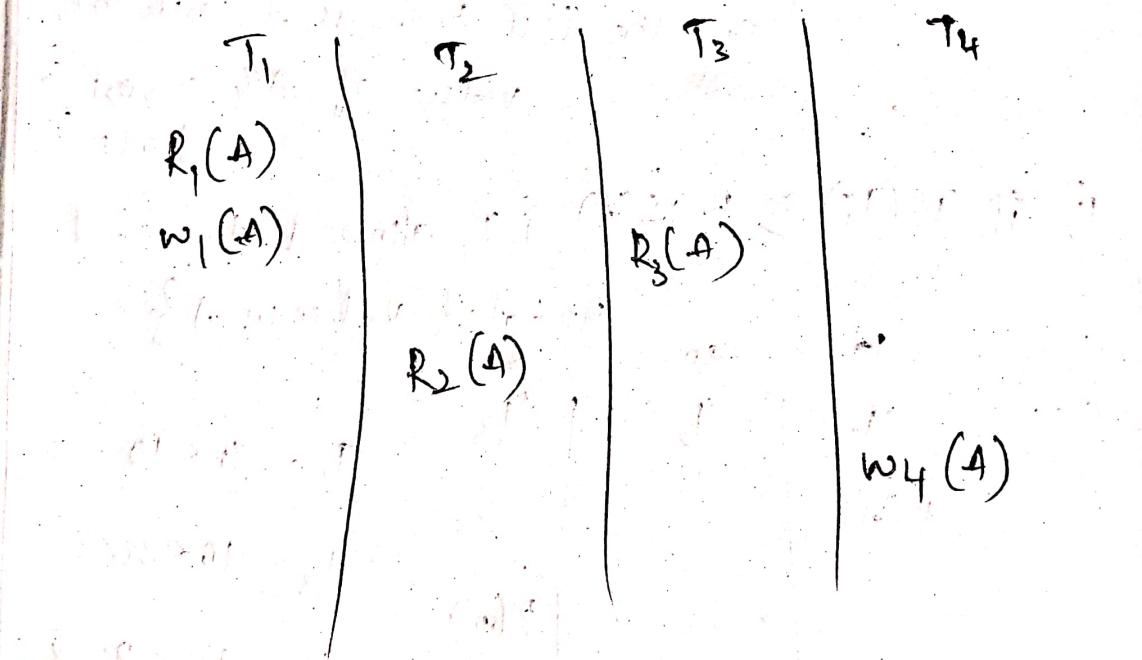
$$T_1 \rightarrow 7:00 \quad Ts(1) < Ts(2)$$

$$T_2 \rightarrow 7:10$$

Rules

- $w - Ts(x)$ is the largest timestamp of any trans that executes writer(x) successfully.
- $r - Ts(x)$ is the largest timestamp of any trans that executes read(x) successfully.

Eg: $12 \quad 25 \quad 36 \quad 42$



$$RTS \Rightarrow RTC = 0$$

$$R_1(A) - T_1 = 12 = 0 < 12 \checkmark$$

$$R_3(A) - T_3 = 36 = 12 < 36 \checkmark$$

$$R_2(A) - T_2 = 25 = 36 < 25 \times \text{stop the transaction}$$

$$RTS = 36$$

Write time stamp

WTS

$$WTS = 0$$

$$w_1(A) = T_1 = 12 \Rightarrow 0 < 12 \checkmark$$

$$w_4(A) = T_4 = 12 \Rightarrow 12 < 42 \times \text{ - stop the trans}$$

$$WTS = 42$$

- ① T_1 issues write (A)

(i) if $TS(T_i) \geq w(A) \& TS(T_i) \geq R(A)$

(ii) if $TS(T_i) < R(A) \times$

(iii) $R(A) \leq TS(T_i) \leq w(A) \times$

Validation Based Protocol

It is an optimistic concurrency control techniques.

In this the transaction is executed in 3 phases

① Read phase-

The T is read & executed. Read the value of various data items and stores them in temporary local variable.

② Validation phase - temporary Var value will be validated against the actual data

to see its serializability.

⑧ Write phase

If the validation of trans is validated,
the temporary results are written to db (or) sys.
Problem - roll back.

start(T) - It contains the time when T starts

validation(T) - It contains the time when
T finishes its read phase & starts validation

finish(T) - It contains the time when T
finishes its write phase.

TS(T) = Validation(T)

↓
Serializability.

All trans starts earlier
must have finish
before transaction T
started i.e
 $\text{finish}(S) < \text{start}(T)$

	<u>T₁</u>	<u>T₂</u>	
<u>Read(B)</u>			<u>S R V W</u>
			<u>T R V W</u>
<u>Read(B)</u>	Read(B)		
	B := B - 50		T ₁ - validation lag
<u>Read(A)</u>		Read(A)	done before T ₂
		A := A + 50	local variable
<u>Read(A)</u>			- No Clobbering
<u><validate></u>			roll back in validation
<u>display(A+B)</u>	<Validates		
	write(A)		
	write(B)		

multiple Granularity:

It is the size of data item allowed

to lock.

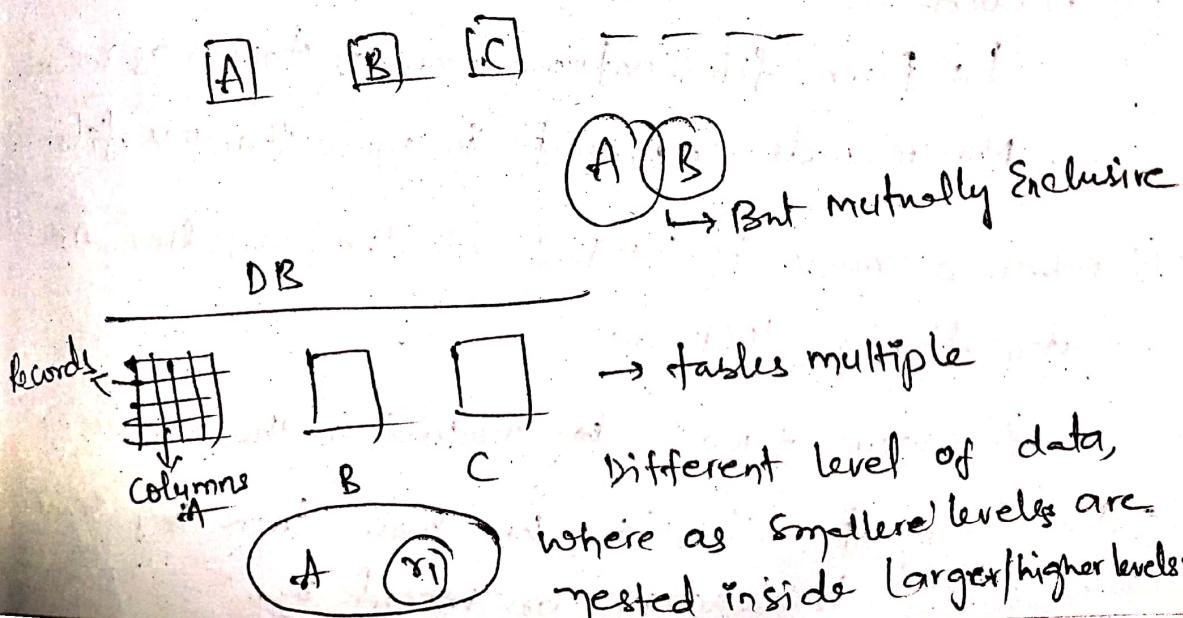
It can be defined as hierarchically breaking up the db into blocks which can be locked. This protocol enhances concurrency and reduces lock overhead.

- It maintains the track of what to lock and how to lock.
- It makes easy to decide either to lock a data item; or to unlock a data item.

This type of hierarchy can be graphically represented as a tree.

We are having each database item as (each) individual data item.

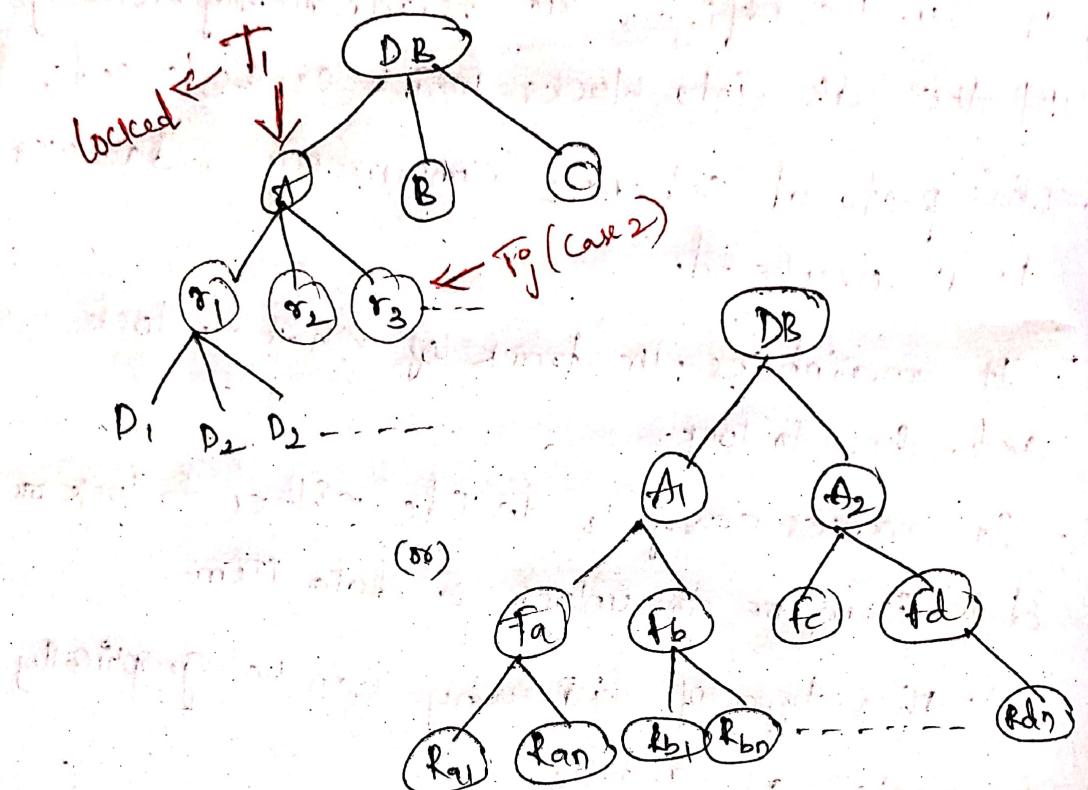
Means there is no intersection b/w d.i.



- we have to use this granularity to ensure Concurrency. (To improve Efficiency)

(i) Data items of multiple sizes

(ii) represented in form of trees



① Database - entire db

② Area - Represents a node of type area.

③ file - children nodes. No file can be present in more than one area.

④ Record

↳ Each file contains nodes known as records.

No records represent in more than one file.

(iii) when a node is locked all the children are automatically locked.

Problems faced by following these rules

- In Emergent situation we have lock by some transaction on record r1.

But σ_1 is already having lock. (T_1)

T_1 wants to acquire another lock on d_2 .

No upper level in this d_2 is ~~record~~ locked.

Case - 1 -

If we have to lock node d_2 , we need to start from root to that node. And we need to check if any transaction will ^{access} be the nodes or not from d_2 .

Case 2 - we have to acquire a lock on DB. we need to check any of the decendent should not be locked by the nearer trans.



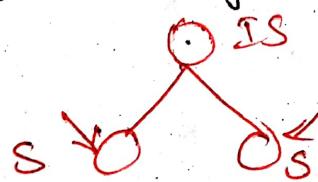
XL - Exclusive
SL - Shared.

Apart from these locks we are using 3 additional lock modes with multiple granularity.

Intention Mode lock

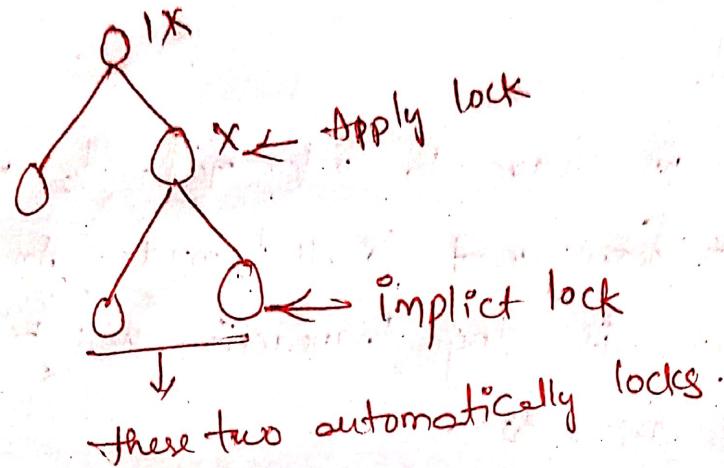
(i) Intention- Shared (IS) :-

It contains explicit locking at a lower level of the tree but only with shared locks.



⑧ Intention Exclusive (IX)

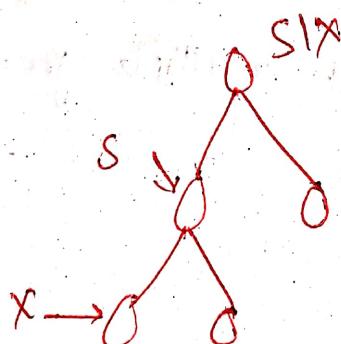
explicit locking at lower level with
exclusive or shared lock. (Atleast one exclusive
lock (X))



⑨ Shared & Intension Exclusive (SIX)

Subtree rooted by that node is locked
explicitly in shared mode & explicit locking
is done at a lower level with exclusive mode.

the node is locked in shared mode,
and some node is locked in exclusive mode
by the same transaction.

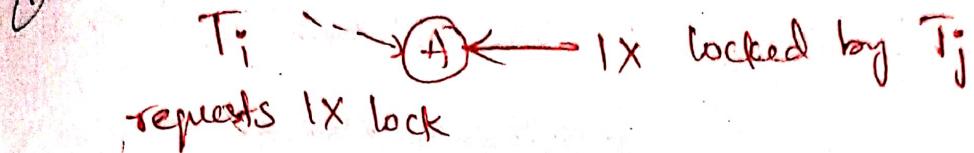


Compatibility Matrix:-

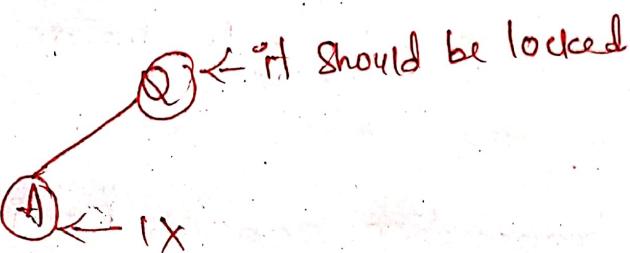
	IS	IX	S	SIX	X
IS	✓	✓	✓	✓	X
IX	✓	✓	X	X	X
S	✓	X	✓	X	X
SIX	✓	X	X	X	X
X	X	X	X	X	X

If a transaction that attempts to lock a node must follow these rules:

① must observe the lock compatibility matrix



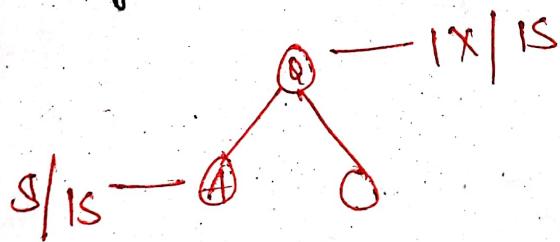
② root of the tree must be locked first and it can be locked in any mode



③ a node 'A' can be locked by T_i in S or

IS mode only if the parent of A is

IS mode only if the parent of A is currently locked by T_i in either IX or IS mode



④ If T_i has not previously unlocked any node only, then the T_i can lock a node.

⑤ If T_i currently has none of the children

of the node-locked only, T_i will unlock a node.

→ Implementation of Atomicity and Durability

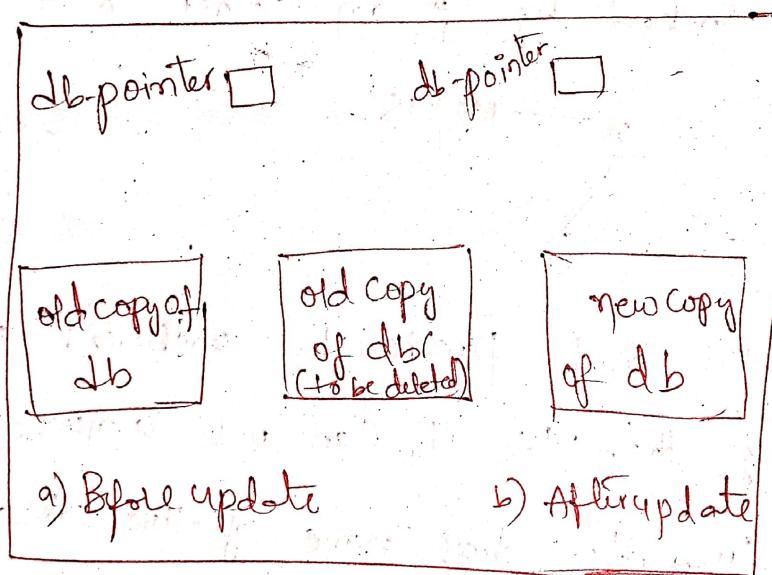
- The recovery management component of a database system implements the support for atomicity and durability.

e.g.: the shadow-database scheme

- All updates are made on a shadow copy of the db.

- db pointer is made to point to the updated shadow copy after.

- the transaction reaches partial commit and
 - All updated pages have been flushed to disk.



- It assumes that only one trans is active at a time.
- Assumes disks do not fail
- Does not handle current transactions

Recover and Atomicity

The purpose of db recovery is to bring the db into last consistent state, which existed before the failure.

To recover from transaction failure, atomicity of transactions as a whole must be maintained - that is either all the operations are executed or none.

There are two types of techniques, which can help a DBMS, in. recovery and maintaining the atomicity of a transaction.

① Log Based Recovery

② shadow paging

Log Based Recovery

The log is a sequence of records which maintains the records of all the update activities performed by a transaction.

It is maintained during the normal transaction processing. It contains enough information to recover from failure.

- log is -the most commonly used structure for recording database.

Types of log records that are maintained in a log:

update log has following operations

① start record (Transaction identifier)

The log record $[T_i, \text{start}]$ is used to indicate that the transaction T_i has started.

② update log Record (Data item identifier)

The log record $[T_i, x, v_1, v_2]$ is used to indicate that the transaction T_i has performed an updated operation on the data item x , having old value v_1 & new value v_2 .

③ read Record - (old value (prior to write))

The log record $[T_i, x]$ is used to indicate that the transaction T_i has read the data item x from the db.

Commit record (After write)

④ Commit record [Ti, Commit] is used to indicate that the transaction Ti has committed.

Abort record

The log record [Ti, abort] is used to indicate that the transaction Ti has aborted & needs to be rollback.

eg	T ₁	T ₂	log record
$A=1000$	read(A)	read(C)	[T ₁ , start]
$B=2000$	$A=A-50$	$C=C-100$	[T ₁ , A]
$C=700$	write(A)	write(C)	[T ₁ , A, 1000, 950]
	read(B)		[T ₁ , B]
	$B=B+50$		[T ₁ , B, 2000, 2050]
	write(B)		[T ₁ , commit]
			[T ₂ , start]
			[T ₂ , C]
			[T ₂ , C, 700 - 600]
			[T ₂ , commit]

② write-ahead log strategy

- log is written before any update is made to db.
- Transaction is not allowed to modify physical db until UNDO is written to stable storage.

operations in Recovery procedure

① UNDO - restore old value

② REDO - updates by new value.

Techniques

① Deferred Db modification:-



Put off to a later time, postpone

It ensures transaction (ensured atomicity) by recording all db modification to the log, but deferring all write operations of a transaction, until the transaction partially commits (i.e. once the final action of the transaction has been executed).

- If the system crashes or if the trans aborts, then the info in the log is ignored.

	<u>T₁</u>	<u>T₂</u>	<u>Case - I</u>
$A = 1000$	read(A)	read(c)	$[T_1, \text{start}]$
$A = A - 50$		$C = C - 100$	$[T_1, A, 950]$ $\xrightarrow{\text{new value}}$
$B = 2000$	write(A)	write(C)	$[T_1, B, 2050]$
$C = 100$	read(B)		$[T_1, \text{Commit}]$
$B = B + 50$			
	write(B)		
		db A [950] B [2050]	failure redo(T_1)
			$[T_1, \text{start} \& T_1, \text{commit present in log}]$
		recovery sum	
			A [950] B [2050]

Case - II

$[T_1, \text{start}]$

$[T_1, A, 950]$

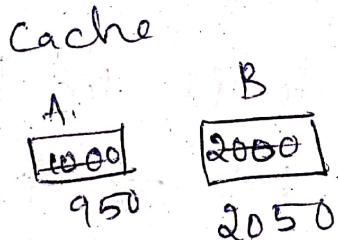
$[T_1, B, 2050]$

failure

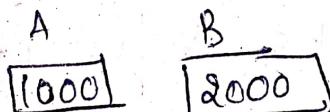
recovery system

no redate (T_1)

as no commit.



database



Case - III

$[T_1, \text{start}]$

$[T_1, A, 750]$

$[T_1, B, 2050]$

$[T_1, \text{Commit}]$

$\rightarrow \text{redo}(T_1)$

$[T_2, \text{start}]$

$[T_2, C, 600]$

failure

Handling the failures, recovery scheme

uses two procedures

① undo(T_i) - It restores the value of data

items updated by trans T_i to the old value.

- T_i needs to undo if the log record

contains T_i start but no T_i commit.

② redo(T_i) -

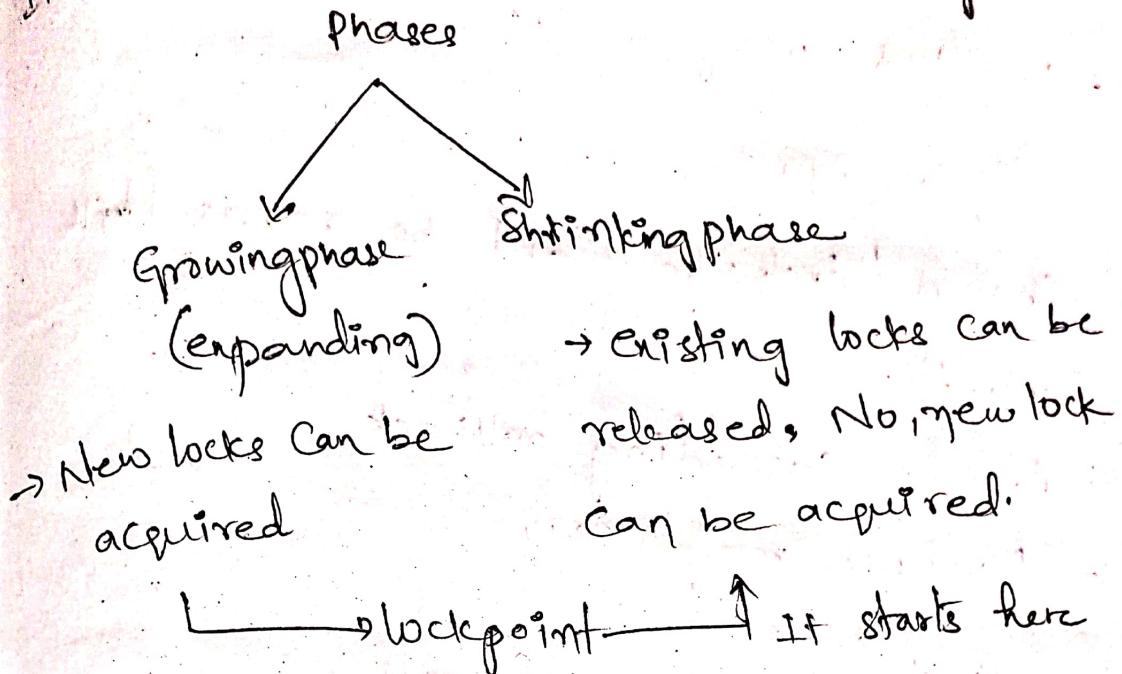
It sets the value of all data items updated by trans T_i to the new value.

- T_i needs to be redo if log record

contains both T_i start & T_i commit

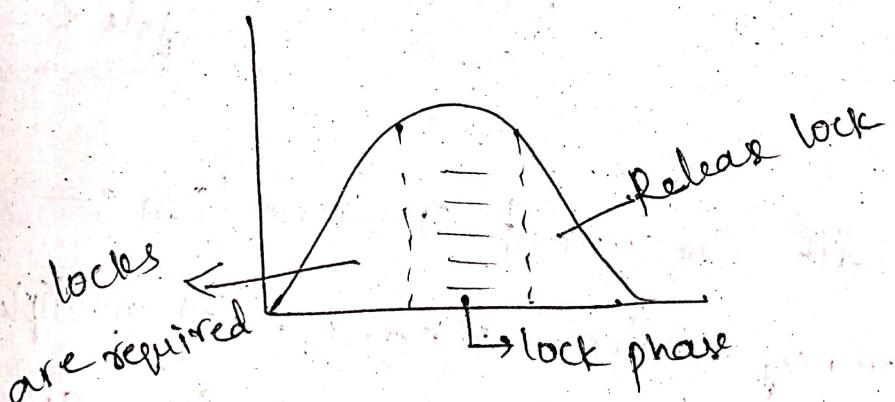
Two phase locking protocol (2PL)

According to the rules of this protocol every trans can be divided into two phases. It requires both locks & unlock being done.



after completion

lock point - It is a point at which transaction has acquired final lock.



Variations of 2PL

- Acquires all lock before it starts
- Release all lock after commit
- AVOIDS cascading rollback

- no deadlock.

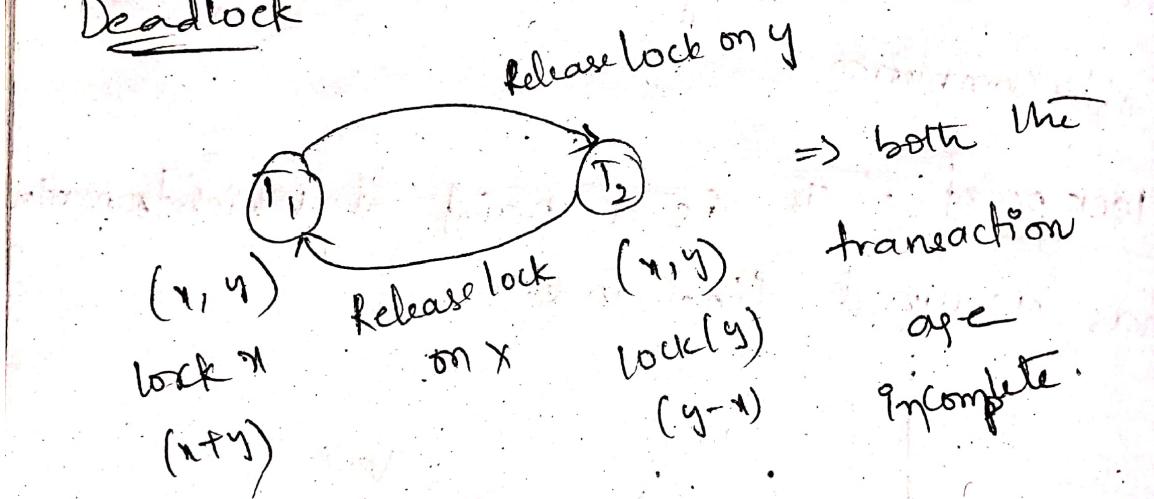
⇒ Strict 2PL

- exclusive lock can't be released until commit.
- helps in cascades & schedule
- deadlock may occur.

⇒ FIFO 2PL

- shared / exclusive can't be released till commit.
- avoid cascading rollback
- deadlock may occur.

Deadlock



Recovery with Concurrent Transaction :-

Concurrency control means that multiple transactions can be executed at the same time and then the interleaved logs occur. But there may be changes in trans results so maintain the order of execution.

of those transactions.

During recovery, it would be very difficult for the recovery system to backtrace all the logs and then start recovering.

Recovery with concurrent transactions can be done in the following four ways.

① Interaction with Concurrency Control

In this, recovery scheme depends greatly on the concurrency control scheme that is used. So, to rollback a failed trans., we must undo the update performed by the transaction.

② Transaction rollback-

- In this scheme, we rollback a failed transaction by using the log.
- The system scans the log backward for every log record of a failed transaction, found in the log the system restores the data item.

③ Checkpoints-

- When more than one transaction is being

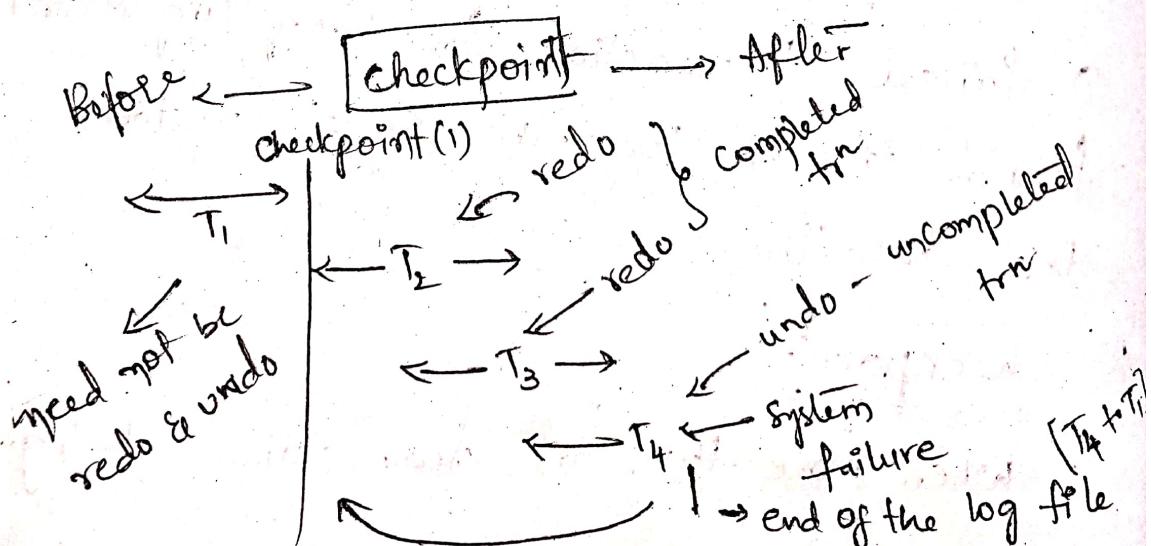
executed then the logs may occur.

During recovery it would become difficult. To make it easy we use 'checkpoint' concept.

The checkpoint is a process of saving a snapshot of the application's state so that it can restart from that point in case of failure.

- It is a type of mechanism where all the previous logs are removed from the system and permanently stored in the storage disk.
- When it reaches the CP, then the trans will be updated into the db, and till that point, the entire log file will be removed from the file. Then the log file is updated with the new step of trans till the next CP & soon.

It is used to declare a point before which the DBMS was in a consistent state, and all the trans were committed.



eg In log file T_2 & T_3 will have $\langle T_m, start \rangle$
8 $\langle T_m, commit \rangle$ T_1 will have only $\langle T_m, commit \rangle$
in log file. Hence it puts T_1, T_2, T_3 in
redo list.

The transaction is put into undo state if the
log is with $\langle T_m, start \rangle$ but no commit or
abort. In this case all the transactions
are undone & logs are removed.

eg: T_4 will have $\langle T_m, start \rangle$ but ~~no~~
~~commit~~ It will be a failed transaction.

(a) Restart Recovery

- When the system recovers from a crash,
it constructs two lists.
- The undo-list consists of transaction to be undone.
- The redo-list consists of transaction to be redone.
- The system constructs two lists.
Initially they are both empty; The system
scans the log backward, examining each record,
until it finds the first \langle checkpoint \rangle record.

ARIES -

Stands for 'Algorithm for recovery and Isolation exploiting Semantics.'

It uses logs to record the progress of transactions and their actions which cause the change to data. It is one of the algorithms in transaction.

It is based on write-ahead log protocol (WAL).

It involves 3 phases

① Analysis

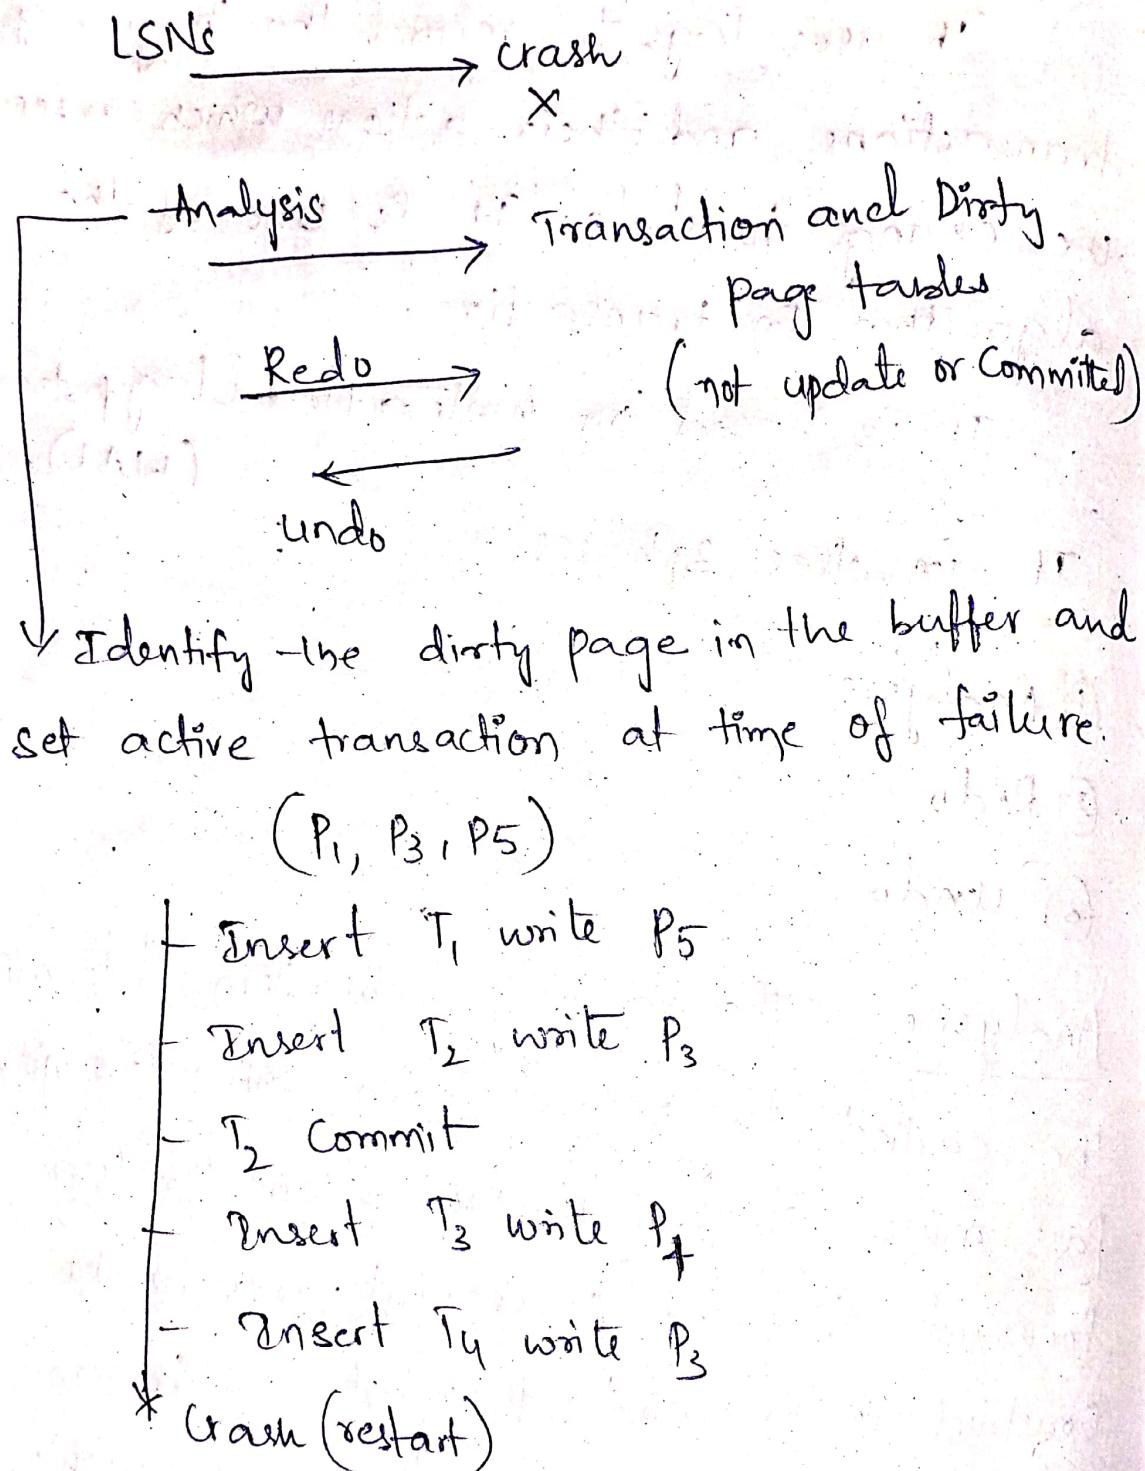
② Redo

③ Undo

Analysis - The recovery subsystem determines the earliest log record from which the next pass must start. It also scans the log forward from the checkpoint record to construct a snapshot of what the system looked like at the instant of the crash.

Redo - Starting at the earliest log sequence number, the log is read forward and each update redone.

Undo - the log is scanned backward and updates corresponding to loser transactions are undone.



When system is restarted analysis is done and it identifies that T₁ & T₃ are active hence restart it.

Redo - find all operations done by DMS before crash and restore the system back to the same state before crash.

It aborts all active transactions those are still in active state at time of crash.

Undo - Scan the log backwards of active transactions in reverse.

T_2 write P_3 in undo

T_3 write P_1

T_1 write P_5